

JSFX SoX Plugins

*A Reimplementation of the SoX Commandline
Processor as JSFX DAW Plugins*

Author: Dr. Thomas Tensi
Date: 2025-03-18
Version: 0.96

Contents

1	Introduction	4
1.1	Overview	4
1.2	Available Effects	5
1.3	Acknowledgements	7
2	Installation of the JSFX-SoX Effects	8
3	Description of the Effects in JSFX-SoX	9
3.1	General Remarks	9
3.2	SoX Allpass Filter	12
3.3	SoX Band Filter	13
3.4	SoX Bandpass Filter	14
3.5	SoX Bandreject Filter	15
3.6	SoX Bass Filter	16
3.7	SoX Biquad Filter	17
3.8	SoX Chorus	18
3.9	SoX Companad	21
3.10	SoX Echo	24
3.11	SoX Echos	26
3.12	SoX Equalizer Filter	28
3.13	SoX Gain	29
3.14	SoX Highpass Filter	30
3.15	SoX Lowpass Filter	31
3.16	SoX MCompanad	32
3.17	SoX Overdrive	34
3.18	SoX Phaser	35

CONTENTS

3.19 SoX Reverb	37
3.20 SoX Treble Filter	38
3.21 SoX Tremolo	39
3.22 Timelocking	40
4 Additional Tools	41
4.1 Oscilloscope	41
5 Regression Test	44
Bibliography	46

Chapter 1

Introduction

1.1 Overview

The JSFX-SoX software package provides plugins for being used in digital audio workstations (DAWs); they implement some of the more prominent audio processing effects from SoX as DAW audio plugins. They implement some of the audio processing effects from SoX. It is also possible to use them in any VST compatible DAW using the freely available ReaJS VST plugin [REAPLUGS] that is able to interpret JSFX files.

SoX [SoXDOC] is a command line audio processing tool for Unix, Windows and MacOSX that transforms source audio files in several formats into other audio files. It provides several standard audio effects (like e.g. filters or reverb) in good quality and with a transparent, open-source implementation. Currently a project called SoX_ng [SoXNG] and led by Martin W. Guy is modernizing all those effects and hence serves as the reference for this JSFX-SoX project.

The effects provided here are a complete rewrite in JSFX of those SoX algorithms nevertheless aiming at producing (*bit-exact*) *identical* renderings in the DAW. This goal is achieved: when rendering some audio externally via SoX and internally with the plugins, there is almost no difference. After rendering and subtracting the results (see chapter 5) those cancel out with typically a residual noise of less than -140dBFS due to rounding or precision errors (SoX often uses 32bit integer processing, while JSFX-SoX always uses double precision floating point processing).

The main motivation for this package is to be able to play around with effects in a DAW and be sure that any external rendering by SoX will produce exactly the same results. Although SoX does not always provide the "best" effects, it still is a reliable and well-defined audio tool.

Only a selection of SoX effects has been reimplemented as plugins, but those are the ones that are — in my opinion — the more prominent effects in that

suite.

Because SoX very often uses rich command line options for its effects, not every effect configuration from SoX can be fully transported into the slider oriented GUI of the JSFX-SoX. For example, the compander of SoX allows the definition of a transfer function having multiple segments. Although the internal engine of the JSFX-SoX compander implements exactly the same internal segment logic of SoX, the user interface only allows the typical definition of a threshold and a compression ratio (leading to a transfer function with three segments).

Note also that a spiffy user interface is *not at all* a priority for this package. Also the parameter ranges are somewhat debatable, but they simply reflect the wide parameter ranges of the SoX command-line effect.

The redesign and restructuring has also been done for easier maintenance, because there is some redundancy and unnecessary complexity in the original sources due to their several contributors. Nevertheless — as pointed out — the effects provided here faithfully model the SoX command-line processing.

All the code is open-source; hence you can check and adapt it to your needs (see chapter ??).

1.2 Available Effects

The following effects are supported:

allpass:

a biquad allpass filter two-poled with filter frequency and the filter bandwith (in several units)

band:

a biquad band filter with center filter frequency and the filter bandwith (in several units) and an option for unpitched audio

bandpass:

a biquad filter for bandpass with center filter frequency and the filter bandwith (in several units)

bandreject:

a biquad filter for bandreject with center filter frequency and the filter bandwith (in several units)

bass:

a biquad filter for boosting or cutting bass with a shelving characteristics with settings for filter frequency and the filter bandwith (in several units)

1.2. AVAILABLE EFFECTS

biquad:

a generic biquad (iir) filter with 6 coefficients b0, b1, b2, a0, a1 and a2

chorus:

a chorus effect with multiple parallel echos modulated either by sine or triangle signals,

compand:

a compander with attack, decay, input gain shift, threshold and compression and soft knee; this is a reduced version of SoX compand with only a simple transfer function and a combined attack/decay setting

echo:

a tapped delay with several absolute delay times and signal decays,

echos:

a sequential delay with delay stages with absolute delay times and signal decays fed with signals from the previous stages,

equalizer:

a biquad filter for equalizing with settings for the pole count, the filter frequency and the filter bandwith (in several units)

gain:

a volume changer by *exact* decibels...

highpass:

a biquad filter for highpass with settings for the pole count, the filter frequency and the filter bandwith (in several units)

mcompand:

a multiband compander with a Linkwitz-Riley crossover filter and for each band a compander with attack, decay, input gain shift, threshold and compression and soft knee; again the companders only allow a simple transfer function and a combined attack/decay setting

lowpass:

a biquad filter for lowpass with settings for the pole count, the filter frequency and the filter bandwith (in several units)

overdrive:

a simple tanh distortion with gain and colour specification

phaser:

a phaser effect with sine or triangle modulation

reverb:

a reverb effect (based on Freeverb) with several parameters for the room (like size and HF damping) as well as a possible predelay

treble:

a biquad filter for boosting or cutting treble with a shelving characteristics with settings for filter frequency and the filter bandwidth (in several units)

tremolo:

a tremolo effect with sine modulation using a double-sideband suppressed carrier modulation

As an analysis tool (e.g. for comparing the tool outputs with reference audio clips) a multilane oscilloscope is also available.

1.3 Acknowledgements

This project is a derivative work based on the foundations laid by the SoX community. Although the algorithms used were modified and redesigned, this project would have been much more complicated and tedious without this basis.

Hence my thanks go to Chris Bagwell, Nick Bailey, Martin W. Guy, Daniel Pouzzner, Måns Rullgård, Rob Sewell and all the other contributors of the SoX project: without your effort this would not have been possible!

Chapter 2

Installation of the JSFX-SoX Effects

The installation is as follows:

1. Load the JSFX-SoX files from the repository in [JSFXSoX].
2. Close the Reaper application or your other DAW using ReaJS (if open).
3. Make a subdirectory `Dr_TT\JSFXSoX` in either the `Effects` directory of the Reaper installation (typically in `\Program Files\Reaper\Effects`) or — if your DAW uses some different directory for JSFX files — in that directory.
4. Copy over all `*.jsfx` and `*.jsfx-inc` files from the distribution `src` directory into the target directory. If helpful, also add the documentation from the root directory.
5. If helpful, also copy the test files from `test` to the effects sub-directory `Dr_TT\JSFXSoX\Test` (see section 5).
6. Restart Reaper (or your other DAW). You should now be able to select the plugins from the JSFX folder (they are all prefixed with "SoX").

Alternatively you can use the ReaPack plugin [REAPACK] and do an automatic install via the `index.xml` file in the repository [JSFXSoX].

Chapter 3

Description of the Effects in JSFX-SoX

3.1 General Remarks

As mentioned in the introduction this package provides several audio tools written in JSFX for emulating SoX bit-exactly.

This goal is reached up to a certain precision (of about -140dBFS), because SoX often uses 32bit integer processing while JSFX-SoX use double precision float processing.

Where noted in the following description, some simplifications have been done to take care of the limited user interface and also some parameters were omitted.

Note again that the focus of this toolset is the faithful reimplementation and somehow a redesign of SoX; a spiffy user interface is *not at all* a priority in this project.

For the same reasons none of the effects of JSFX-SoX displays anything; they just process audio parametrized by their slider settings.

Nevertheless for debugging purposes, some textual display is supported via keypresses in the effects window as follows:

- Internal effect data and debug instrumentation can be displayed by pressing “d”(ebug). Figure 3.2 shows the debug data for a multiband-compander as an example.
- The table of allocated strings is shown via key “s”(trings).
- The local memory is displayed via “m”(emory) (see figure 3.1).
- All those displays scale automatically and are scrollable via “+” and

3.1. GENERAL REMARKS

SoX Reverb							Edit...
is wet only? WetAndDry							▼
		Reverberance [%]					60.0
		HF Damping [%]	<input type="checkbox"/>				22.0
		Room Scale [%]			<input type="checkbox"/>		87.5
		Stereo Depth [%]	<input type="checkbox"/>				34.88
		Pre Delay [ms]	<input type="checkbox"/>				20.0
		Wet Gain [dB]	<input type="checkbox"/>				-3.0
1485:	0	1508:	0	1531:	+1.2010e-003	1554:	882 1577: 0
1486:	589	1509:	0	1532:	+5.6183e-004	1555:	3 1578: 0
1487:	0	1510:	2	1533:	2	1556:	0 1579: 0
1488:	590	1511:	-2.2176e-004	1534:	-1.6445e-003	1557:	0 1580: 0
1489:	0	1512:	+5.8978e-004	1535:	+6.1773e-004	1558:	0 1581: 0
1490:	591	1513:	8	1536:	8	1559:	0 1582: 0
1491:	0	1514:	1116	1537:	0	1560:	0 1583: 0
1492:	592	1515:	1188	1538:	+9.1708e-001	1561:	0 1584: 0
1493:	0	1516:	1277	1539:	+2.6600e-001	1562:	0 1585: 0
1494:	593	1517:	1356	1540:	+2.0000e-002	1563:	0 1586: 0
1495:	0	1518:	1422	1541:	+3.4880e-001	1564:	0 1587: 0
1496:	594	1519:	1491	1542:	+1.0619e-002	1565:	0 1588: 0
1497:	0	1520:	1557	1543:	+8.8750e-001	1566:	0 1589: 0
1498:	595	1521:	1617	1544:	1546	1567:	0 1590: 0
1499:	0	1522:	4	1545:	2	1568:	0 1591: 0
1500:	596	1523:	225	1546:	1549	1569:	0 1592: 0
1501:	0	1524:	341	1547:	49298	1570:	0 1593: 0
1502:	597	1525:	441	1548:	3	1571:	0 1594: 0
1503:	0	1526:	556	1549:	1553	1572:	0 1595: 0
1504:	598	1527:	2	1550:	2	1573:	0 1596: 0
1505:	0	1528:	1531	1551:	23607	1574:	0 1597: 0
1506:	599	1529:	1534	1552:	22053	1575:	0 1598: 0
1507:	2	1530:	2	1553:	22050	1576:	0 1599: 0

Figure 3.1: Memory View for JSFX-SoX Effects in Reaper

“-” and reset via “0” (with “*” and “_” allow for scrolling by a larger increment).

- The font size in those displays can be scaled up or down by “3” and “1” and reset via “2”.
- The display of the effect (if any) is shown via “e”(ffect) which is the default (and — as mentioned before — normally empty).

All plugins are designed such that view are kept visible even when playback state is changed. Hence one can e.g. have a live look at the memory data for some plugin (provided you know its layout in memory).

All effects of JSFX-SoX are discussed in alphabetical order in the following chapter. Note that the effects description is mostly taken from the SoX documentation [SoXDOC, SoXNG] except for specifics of the JSFX-SoX effects.

CHAPTER 3. DESCRIPTION OF THE EFFECTS IN JSFX-SOX




Figure 3.2: Debug View for JSFX-SoX Effects in Reaper

3.2 SoX Allpass Filter




Figure 3.3: Panel for SoX Plugin “Allpass”

Parameter	Description	Unit
Frequency	the center frequency of the filter	Hz
Bandwidth	the bandwidth modulus of the filter	—
Bandwidth Unit	the bandwidth unit of the filter	Frequency / Octaves / Quality / Butterworth

This effect is a two-pole all-pass filter with **Frequency** as center frequency and filter-width **Bandwidth** with unit **Bandwidth Unit**. The bandwidth kinds are a relative *Frequency*, a specification of *octaves*, the filter *quality* or the *butterworth* quality (with fixed quality $q = \sqrt{2}/2$).

An all-pass filter changes the audio’s frequency-to-phase relationship without changing its frequency-to-amplitude relationship. The detailed filter description can be found in [RBJFILT].

3.3 SoX Band Filter




Figure 3.4: Panel for SoX Plugin “Band”

Parameter	Description	Unit
Unpitched Mode?	flag to tell whether special processing for unpitched audio is applied	Boolean
Frequency	the center frequency of the filter	Hz
Bandwidth	the bandwidth modulus of the filter	—
Bandwidth Unit	the bandwidth unit of the filter	Frequency / Octaves / Quality / Butterworth

This effect is a band-pass filter. The frequency response drops logarithmically around **Frequency**, the **Bandwidth** and **Bandwidth Unit** parameters gives the slope of the drop; frequencies at *frequency+width* and *frequency-width* will be half of their original amplitudes. The effect defaults to a mode that is oriented to pitched audio, i.e. voice, singing, or instrumental music.

When the option **Unpitched Mode?** is set, an alternate mode for un-pitched audio (e.g. percussion) is applied. Note that this option introduces a power-gain of about 11dB in the filter, so beware of output clipping; the option introduces noise in the shape of the filter, i.e. peaking at the center frequency and settling around it.

3.4 SoX Bandpass Filter




Figure 3.5: Panel for SoX Plugin “Bandpass”

Parameter	Description	Unit
Mode	the kind of the filter (here: BandPass)	BandPass / BandReject
Cst. Skirt Gain?	flag to tell whether a constant skirt gain is applied	Boolean
Frequency	the center frequency of the filter	Hz
Bandwidth	the bandwidth modulus of the filter	—
Bandwidth Unit	the bandwidth unit of the filter	Frequency / Octaves / Quality / Butterworth / Slope

This effect is a bandpass filter; by selecting the appropriate **Filter Kind**, this effect is a two-pole Butterworth band-pass filter with **Frequency** as central frequency, and (3dB-point) band-width given by **Bandwidth** and **Bandwidth Unit**. The **Cst. Skirt Gain?** option selects a constant skirt gain (peak gain = Q) instead of the default, which is a constant 0dB peak gain. The filters roll off at 6dB per octave (20dB per decade).

The detailed filter description can be found in [RBJFILT].

3.5 SoX Bandreject Filter




Figure 3.6: Panel for SoX Plugin “Bandreject”

Parameter	Description	Unit
Mode	the kind of the filter (here: BandReject)	BandPass / BandReject
Frequency	the center frequency of the filter	Hz
Bandwidth	the bandwidth modulus of the filter	—
Bandwidth Unit	the bandwidth unit of the filter	Frequency / Octaves / Quality / Butterworth / Slope

This effect is a bandreject filter; by selecting the appropriate **Filter Kind**, this effect is a two-pole Butterworth band-reject filter with **Frequency** as central frequency, and (3dB-point) band-width given by **Bandwidth** and **Bandwidth Unit**.

The detailed filter description can be found in [RBJFILT].

3.6 SoX Bass Filter




Figure 3.7: Panel for SoX Plugin “Bass”

Parameter	Description	Unit
Mode	the kind of the filter (here: Bass)	Bass / Treble
Gain	gain of filter at 0Hz	dB
Frequency	the center frequency of the filter	Hz
Bandwidth	the bandwidth modulus of the filter	—
Bandwidth Unit	the bandwidth unit of the filter	Frequency / Octaves / Quality / Butterworth

This effect is a bass filter; by selecting the appropriate **Filter Kind**, this effect boosts or cuts the bass (lower) frequencies of the audio using a two-pole shelving filter with a response similar to that of a standard hi-fi’s tone-controls. This is also known as shelving equalisation (EQ).

The parameters are as follows:

- **Gain** gives the gain at 0Hz. Its useful range is about -20 (for a large cut) to +20 (for a large boost). Beware of clipping when using a positive gain.
- **Frequency** sets the filter’s central frequency and can be used to extend or reduce the frequency range to be boosted or cut.
- The band-width given by parameters **Bandwidth** and **Bandwidth Unit** determines how steep is the filter’s shelf transition. In addition to the common width specification methods described above, “slope” may be used. The useful range of slope is about 0.3, for a gentle slope, to 1 (the maximum), for a steep slope.

The detailed filter description can be found in [RBJFILT].

3.7 SoX Biquad Filter




Figure 3.8: Panel for SoX Plugin “Biquad”

Parameter	Description	Unit
b0	coefficient for x_n	—
b1	coefficient for x_{n-1}	—
b2	coefficient for x_{n-2}	—
a0	coefficient for y_n	—
a1	coefficient for y_{n-1}	—
a2	coefficient for y_{n-2}	—

This effect is a biquad IIR filter with the given coefficients (see [DBIQFILT]). It implements the (direct form) function

$$y_n = \sum_{i=0}^2 b_i x_{n-i} - \sum_{i=1}^2 a_i y_{n-i}$$

and is the basis for the other biquad filters (like e.g. the “SoX Equalizer”).

The signal flow graph of the biquad filter is given in figure 3.9 and implements the above equation.




Figure 3.9: Signal Flow Graph for SoX Plugin “Biquad”

3.8. SOX CHORUS

3.8 SoX Chorus




Figure 3.10: Panel for SoX Plugin “Chorus”

Parameter	Description	Unit
Effect Kind	the kind of the delay (here: Chorus)	Chorus / Delay Sequence / Tapped Delay
In Gain	the gain factor before processing	—
Out Gain	the gain factor applied after processing	—
Time Offset	the point in project time where modulation is at phase 0° (see 3.22)	s
Stage Count	the count of the delay stages	—
Stage Index	the index of the delay stage to be adapted	—
for each stage		
Delay	the <i>absolute</i> delay duration for that stage	ms
Decay	the decay factor for that stage	—
Frequency	the modulation frequency for that stage	Hz
Depth	the modulation depth for that stage	ms
Waveform	the modulation waveform for that stage	Sine / Triangle

This effect is a variant of the plugin `SoxChorusAndEcho` and implements a modulated tapped delay (a “chorus”) for the audio.

In Gain is the amplification factor for the input. **Out Gain** is the amplification factor of the output.

Because this effect uses several modulations there is a parameter **Time Offset** locking those modulators to a given start time i.e. this effect is *time-locked*. For details refer to section 3.22.

Stage Count gives the number of stages in the delay, i.e. pairs of (delay, decay), which can be set to the desired number, where the maximum is 10. **Stage Index** selects the delay stage whose parameters shall be modified; note that it is possible to modify a stage only when its index is not greater than the stage count.

Delay gives the delay in milliseconds for the selected delay stage, **Decay** a factor for the decay of that delay stage. **Frequency** gives the frequency of the delay modulation in Hz and **Depth** specifies the amplitude of the modulation in milliseconds. Possible waveforms are sine or triangle and can be specified by **Waveform Kind**.

The signal flow graph of the chorus is given in figure 3.11: The input signal is amplified by **inGain** and then flows to the aggregate summation as well as to each chorus stage. Each stage delays the signal by delay_i modulated by depth_i with frequency_i and finally amplifies the result by decay_i and also feeds it to the summation. The summation result is amplified by **outGain** and fed to the output of the effect.

Note that there is an (unchecked!) recommendation to have

$$\text{inGain} \times \sum_i \text{decay}_i > \frac{1}{\text{outGain}}$$

3.8. SOX CHORUS




Figure 3.11: Signal Flow Graph for SoX Plugin “Chorus”

3.9 SoX Compand




Figure 3.12: Panel for SoX Plugin “Compand”

Parameter	Description	Unit
Band Count	the count of the bands (here: 1)	—
Band Index	the index of the band to be adapted (here: 1)	—
Attack	the attack time of the compander	s
Decay	the decay time of the compander	s
Knee	the rounding of the corners in the transfer function	dB
Threshold	the threshold of the compander	dBFS
Ratio	the compression factor of the compander	—
Gain	the compander gain before processing	dB
Top Frequency	the compander band top frequency (for all but the last active band)	Hz

This effect implements a compander to compress or expand the dynamic range of the audio. A compander and multiband compander are both variants of the plugin `SoXCompander` where a simple compander is a multiband compander with just one band.

Band Count tells the number of bands in a multiband compander, for a single band compander this must be set to 1. **Band Index** selects the band whose parameters shall be modified, for a single band compander this also must be set to 1, because there is only one active band.

3.9. SOX COMPAND




Figure 3.13: Example Transfer Function (Threshold 20dBFS, Ratio 4:1)

The parameters **Attack** and **Decay** (in seconds) determine the time over which the instantaneous level of the input signal is averaged to determine its volume; attacks refer to increases in volume and decays refer to decreases. For most situations, the attack time (response to the music getting louder) should be shorter than the decay time because the human ear is more sensitive to sudden loud music than sudden soft music. Typical values are 0.3s for attack and 0.8s for decay.

The transfer function of the compander is given by parameters **Threshold**, **Ratio** and **Knee**. The compander leaves the original level unchanged, when it is below threshold and compresses it by ratio beyond this threshold. So e.g. for a threshold of 20dBFS, a knee of 0dB and a ratio of 4:1 the transfer function is a graph shown in figure 3.13. Note that for technical reasons SoX uses a linear lead-in segment with size 10dB below threshold value.

If the parameter **Knee** is greater than 0, the corner points of the transfer function will be rounded by that amount.

The parameter **Gain** is an additional gain in dB to be applied at all points on the transfer function and allows easy adjustment of the overall gain.

As mentioned before the compander is a multiband compander with one band; hence the setting for the **Top Frequency** is not used.

- Restriction:** Only one overall pair of attack/decay parameters may be specified (where SoX allows one pair per channel). This is in principle supported by the effects engine of JSFX-SoX, but not supported in the current user interface.
- Restriction:** The original SoX allows an arbitrary multi-segmented transfer function. This is in principle supported by the effects engine of JSFX-SoX, but not supported in the current user interface.
- Restriction:** There is no delay parameter for delayed compansion.

3.10 SoX Echo




Figure 3.14: Panel for SoX Plugin “Echo”

Parameter	Description	Unit
Effect Kind	the kind of the delay (here: Tapped Delay)	Chorus / Delay Sequence / Tapped Delay
In Gain	the gain factor before processing	—
Out Gain	the gain factor applied after processing	—
Stage Count	the count of the delay stages	—
Stage Index	the index of the delay stage to be adapted	—
for each stage		
Delay	the <i>absolute</i> delay duration for that stage	ms
Decay	the decay factor for that stage	—

This effect is a variant of the plugin `SoxChorusAndEcho` and implements a tapped delay for the audio.

In Gain is the amplification factor for the input. **Out Gain** is the amplification factor of the output.

Stage Count gives the number of stages in the delay, i.e. pairs of (delay, decay), which can be set to the desired number, where the maximum is 10. **Stage Index** selects the delay stage whose parameters shall be modified; note that it is possible to modify a stage only when its index is not greater than the stage count.

Delay gives the delay in milliseconds for the selected delay stage, **Decay** a factor for the decay of that delay stage.

The signal flow graph of the delay in sequence is given in figure 3.15: The input signal is amplified by `inGain` and then flows to the aggregate summation as well as to each stage. Each stage delays the signal by `delayi`, amplifies the

result by decay_i and also feeds it to the summation. The summation result is amplified by outGain and fed to the output of the effect.




Figure 3.15: Signal Flow Graph for SoX Plugin “Echo”

Note that there is an (unchecked!) recommendation to have

$$\text{inGain} \times \sum_i \text{decay}_i > \frac{1}{\text{outGain}}$$

3.11. SOX ECHOS

3.11 SoX Echoes




Figure 3.16: Panel for SoX Plugin “Echos”

Parameter	Description	Unit
Effect Kind	the kind of the delay (here: Delay Sequence)	Chorus / Delay Sequence / Tapped Delay
In Gain	the gain factor applied to the input before summation	—
Out Gain	the gain factor applied after processing	—
Stage Count	the count of the delay stages	—
Stage Index	the index of the stage to be adapted	—
for each stage		
Delay	the <i>absolute</i> delay duration for that stage	ms
Decay	the decay factor for that stage	—

This effect is a variant of the plugin `SoxChorusAndEcho` and implements a sequential delay for the audio.

In Gain is the amplification factor for the input before going into summation (see signal flow graph), **Out Gain** is the amplification factor of the output.

Stage Count gives the number of stages in the delay, i.e. pairs of (delay, decay), which can be set to the desired number, where the maximum is 10. **Stage Index** selects the delay stage whose parameters shall be modified; note that it is possible to modify a stage only when its index is not greater than the stage count.

Delay gives the delay in milliseconds for the selected delay stage, **Decay** a factor for the decay of that delay stage.

The signal flow graph of the delay in sequence is given in figure 3.17: The input signal is amplified by **inGain** and then flows to the aggregate summation and each stage. Each stage delays the signal by **delay_i**, amplifies the result by

decay_i and also feeds it to the summation. The summation result is amplified by `outGain` and fed to the output of the effect.




Figure 3.17: Signal Flow Graph for SoX Plugin “Echos”

Note that there is an (unchecked!) recommendation to have

$$\text{inGain} \times \sum_i \text{decay}_i > \frac{1}{\text{outGain}}$$

3.12 SoX Equalizer Filter




Figure 3.18: Panel for SoX Plugin “Equalizer”

Parameter	Description	Unit
Frequency	the 3dB point frequency of the filter	Hz
Bandwidth	the bandwidth modulus of the filter	—
Bandwidth Unit	the bandwidth unit of the filter	Frequency / Octaves / Quality / Butterworth
Eq. Gain	gain of filter at frequency	dB

This effect is a two-pole peaking equalisation (EQ) filter. With this filter, the signal-level at and around a selected frequency can be increased or decreased, whilst (unlike band-pass and band-reject filters) that at all other frequencies is unchanged.

The parameter **Frequency** gives the filter’s central frequency in Hz, parameters **Bandwidth** and **Bandwidth Unit** the bandwidth and **Gain** the required amplification or attenuation in decibels. Beware of clipping when using a positive gain.

The filter is described in detail in [RBJFILT].

3.13 SoX Gain



Figure 3.19: Panel for SoX Plugin “Gain”

Parameter	Description	Unit
Gain	the amplification or attenuation factor	dB

This effect is an amplifier or attenuator for the audio signal with a single Gain parameter in decibels. The gain factor applies to all channels identically.

Nothing special, but note that the calculation is exact, hence a gain of -6dB does *not* halve the signal (but a gain of -6.0206dB does quite well).

3.14 SoX Highpass Filter




Figure 3.20: Panel for SoX Plugin “Highpass”

Parameter	Description	Unit
Mode	the kind of the filter (here: HighPass)	HighPass / LowPass
Number of Poles	selects between single and double pole filter	single/double
Frequency	the 3dB point frequency of the filter	Hz
Bandwidth	the bandwidth modulus of the filter	—
Bandwidth Unit	the bandwidth unit of the filter	Frequency / Octaves / Quality / Butterworth

This effect is a high-pass filter with a 3dB point **Frequency**. Depending on **Number of Poles** the filter can be either single-pole or double-pole. The parameters **Bandwidth** and **Bandwidth Unit** apply only to double-pole filters; a Butterworth response is given by butterworth selection or by a q of 0.707. The filters roll off at 6dB per pole per octave (20dB per pole per decade).

The double-pole filters are described in detail in [RBJFILT].

3.15 SoX Lowpass Filter




Figure 3.21: Panel for SoX Plugin “Lowpass”

Parameter	Description	Unit
Mode	the kind of the filter (here: LowPass)	HighPass / LowPass
Number of Poles	selects between single and double pole filter	single/double
Frequency	the 3dB point frequency of the filter	Hz
Bandwidth	the bandwidth modulus of the filter	—
Bandwidth Unit	the bandwidth unit of the filter	Frequency / Octaves / Quality / Butterworth

This effect is a low-pass filter with a 3dB point **Frequency**. Depending on **Number of Poles** the filter can be either single-pole or double-pole. The parameters **Bandwidth** and **Bandwidth Unit** apply only to double-pole filters; a Butterworth response is given by butterworth selection or by a q of 0.707. The filters roll off at 6dB per pole per octave (20dB per pole per decade).

The double-pole filters are described in detail in [RBJFILT].

3.16 SoX MCompanad




Figure 3.22: Panel for SoX Plugin “MCompanad”

Parameter	Description	Unit
Band Count	the count of the bands	—
Band Index	the index of the band to be adapted	—
for each band		
Attack	the attack time of the compander band	s
Decay	the decay time of the compander band	s
Knee	the rounding of the corners in the transfer function	dB
Threshold	the threshold of the compander band	dBFS
Ratio	the compression factor of the compander band	—
Gain	the compander band gain before processing	dB
Top Frequency	the compander band top frequency (for all but the last active band)	Hz

Compander and multiband compander are both variants of the plugin **SoX-Compander** where a simple compander is a multiband compander with just one band. For a general multi-band compander the audio is first divided into bands using Linkwitz-Riley cross-over filters and later separately specifiable compander run on every band (see the compand effect in 3.9 for the definition of its parameters).

Band Count gives the number of bands in a multiband compander which can be set to the desired number, where the maximum is 10. **Band Index** selects

the band whose parameters shall be modified; note that it is possible to modify a band only when its index is not greater than the band count.

Restriction: Only one overall pair of attack/decay parameters may be specified per band (where SoX allows one pair per channel). This is in principle supported by the effects engine of JSFX-SoX, but not supported in the current user interface.

Restriction: The original SoX allows an arbitrary multi-segmented transfer function. This is in principle supported by the effects engine of JSFX-SoX, but not supported in the current user interface.

Restriction: There is no delay parameter for delayed compansion.

3.17 SoX Overdrive



Figure 3.23: Panel for SoX Plugin “Overdrive”

Parameter	Description	Unit
Gain	the overdrive gain before processing	dB
Colour	percentage for the amount of even harmonic content in output	—

This effect implements an tanh overdrive. **Gain** gives the input gain in decibels, the parameter **Colour** controls the amount of even harmonic content in the overdriven output.

3.18 SoX Phaser




Figure 3.24: Panel for SoX Plugin “Phaser”

Parameter	Description	Unit
Effect Kind	the kind of the modulation (here: Phaser)	Phaser / Tremolo
In Gain	the gain factor before processing	—
Out Gain	the gain factor applied after processing	—
Delay	the predelay of the effect	ms
Decay	the decay factor of the phaser	—
Frequency	the phaser modulation frequency	Hz
Waveform	the modulation waveform	Sine / Triangle
Time Offset	the point in project time where modulation is at phase 0° (see 3.22)	s

This effect is a phaser effect to the audio.

In Gain is the amplification factor for the input. **Delay** gives the delay in milliseconds, **Decay** a factor for the decay within the phaser and **Frequency** gives the modulation frequency in Hz. The **Waveform** of the modulation is either sinusoidal — preferable for multiple instruments — or triangular — gives single instruments a sharper phasing effect —. **Out Gain** is the amplification factor of the output.

The decay should be less than 0.5 to avoid feedback, and usually no less than 0.1. The (unchecked!) recommendation is to have

$$\begin{aligned} \text{in_gain} &< 1 - \text{decay}^2 \\ \text{out_gain} &< \frac{1 - \text{decay}}{\text{in_gain}} \end{aligned}$$

3.18. SOX PHASER

The parameter **Time Offset** shows that this effect is time-locked. For details refer to section 3.22.

3.19 SoX Reverb




Figure 3.25: Panel for SoX Plugin “Reverb”

Parameter	Description	Unit
Is Wet Only?	tells whether only the wet signal should be produced	Boolean
Reverberance	percentage for reverb density	—
HF Damping	percentage amount of damping of high frequencies for every reflection relative to low frequencies	—
Room Scale	percentage for size of the room (more precisely the reflectivity of the room)	—
Stereo Depth	percentage amount of stereo effect	—
Predelay	time offset until first reverb occurs	ms
Wet Gain	gain of wet signal relative to dry signal	dB

This effect implements reverberation of audio using the “freeverb” algorithm, which uses eight parallel Schröder-Moorer filtered-feedback comb-filters followed by four Schröder allpasses in series.

Details on this algorithm can be found in [FREEVERB].

3.20 SoX Treble Filter




Figure 3.26: Panel for SoX Plugin “Treble”

Parameter	Description	Unit
Mode	the kind of the filter (here: Treble)	Bass / Treble
Gain	gain of filter at 22kHz	dB
Frequency	the center frequency of the filter	Hz
Bandwidth	the bandwidth modulus of the filter	—
Bandwidth Unit	the bandwidth unit of the filter	Frequency / Octaves / Quality / Butterworth

This effect is a treble filter; by selecting the appropriate **Filter Kind**, this effect boosts or cuts the treble (upper) frequencies of the audio using a two-pole shelving filter with a response similar to that of a standard hi-fi’s tone-controls. This is also known as shelving equalisation (EQ).

The parameters are as follows:

- **Gain** gives the gain at a frequency whichever is the lower of 22kHz and the Nyquist frequency. Its useful range is about -20 (for a large cut) to +20 (for a large boost). Beware of clipping when using a positive gain.
- **Frequency** sets the filter’s central frequency and can be used to extend or reduce the frequency range to be boosted or cut.
- The band-width given by parameters **Bandwidth** and **Bandwidth Unit** determines how steep is the filter’s shelf transition. In addition to the common width specification methods described above, “slope” may be used. The useful range of slope is about 0.3, for a gentle slope, to 1 (the maximum), for a steep slope.

The detailed filter description can be found in [RBJFILT].

3.21 SoX Tremolo



Figure 3.27: Panel for SoX Plugin “Tremolo”

Parameter	Description	Unit
Effect Kind	the kind of the modulation (here: Tremolo)	Phaser / Tremolo
Frequency	the modulation frequency of the tremolo	Hz
Depth	percentage value for the intensity of modulation	—
Time Offset	the point in project time where modulation is at phase 0° (see 3.22)	s

This effect is a tremolo effect. This tremolo is done by signal multiplication; hence it is a low frequency double sideband suppressed carrier modulation. Parameter **Frequency** gives the tremolo frequency, **Depth** gives the intensity as a percentage.

Time Offset shows that this effect is time-locked. For details refer to section 3.22.

3.22 Timelocking

There are effects that behave differently in time, technically they are *time-variant*. A filter does not care *when* a signal arrives, but a modulated effect like e.g. a phaser produces a different sound for different start times because the modulation is normally in another phase.

Hence when looking at the behaviour at a specific point in time, those time-variant effects would behave differently when the effect start time were varied.

For example, assume a phaser with a 0.25Hz modulation (one cycle every 4s): when you start the effect 1s later, its modulation is now off by 90°. This is not helpful when the effect now depends on start time or loop positioning.

To circumvent this problem, all time variant effects from above (phaser and tremolo) are *time-locked* i.e. they check the current play position and always behave the same at some specific point in time regardless of the playback start time.

Additionally those effects have a parameter called **Time Offset**. This parameter tells at what time the effect has a phase of zero in its modulation. The default is 0s, but it may be adapted accordingly.

Take the phaser above and assume you want to make sure that its modulation is exactly at 0° at position 155s within your song¹. Then you just set **Time Offset** to “155”. Because the period of the modulation is 4s it is also okay to use $155 + 4k, k \in \mathbb{Z}$ as offset (e.g. “3”), but the above saves you from some calculation for complicated modulation frequencies.

By this method even time-variant effects can be synced with externally generated audio material.

¹This is a little lie, because the initial phaser modulation phase is 90°, but the argument is still valid.

Chapter 4

Additional Tools

For the analysis of the plugins one can use the stock plugins from Reaper, e.g. the spectrum analyzer `gfonalyzer`.

While this plugin is very helpful, unfortunately there is no really good oscilloscope available with parallel lanes and a precise time axis. Hence this is added to the plugin suite.

4.1 Oscilloscope




Figure 4.1: Oscilloscope with Parallel Lanes

4.1. OSCILLOSCOPE

Parameter	Description	Unit
Time Offset	the offset between the start of the recording and the left position of the oscilloscope window	s
Time Range	the displayed time range in the oscilloscope window	s
Zoom Factor	the magnification factor for the signals in the window	dB
Lane Count	the number of parallel lanes displayed	—

The oscilloscope is based on Cockos' gfxscope, but, in contrast, it allows the parallel display of several signal lanes (either 1, 2, 3 or 4) where the number of parallel lanes can be selected via a context menu on the menu button **Lane Count**.

There is a fixed assignment of channels to the lanes: lane 1 displays track channels 1 and 2, lane 2 track channels 3 and 4 and so on. Both stereo channels are overlaid in the lane and encoded with different colors.

All lanes have a vertical grid with an adaptive dB scaling. The displayed signals may be amplified or attenuated by some magnification factor set by the button **Zoom**. The button has a context menu (with a ten decibels resolution), but its value can also be changed by dragging the mouse on that menu button, by a vertical mouse drag within the oscilloscope window or by using the mouse wheel with the SHIFT-key pressed. When the CTRL-key is pressed during dragging, its sensitivity is reduced for a finer control. A double-click on that menu button sets its value to the default of 0dB.

The time resolution of the window can be set by the menu button **Time Range**. The button has a context menu with decadic values, but for a finer setting also a mouse drag can be done on the button. Also a horizontal mouse drag with the SHIFT-key pressed changes that value. A double-click on that menu button sets the value to the default of the maximum available time range (typically about 20s).

The oscilloscope records the signals for several seconds (depending on the sample rate, typically 20s for 44.1kHz) and the internal sample buffer is reset as soon as playback is started. Samples are kept, but subsequently rolled out in a fifo fashion when the buffer is full. Note that the start time of the recording is always displayed as the leftmost position on the horizontal axis of the window (in the unit "samples" on the top margin and in the unit "seconds" on the bottom margin).

To have a look within the recorded signals it is possible to shift the displayed start time relative to the start position by setting the menu button **Time Offset**. This can be done via a context menu of the button (with half a



Figure 4.2: Comparing a Test and a Reference Signal

second resolution), via dragging on the button, by a horizontal mouse drag within the oscilloscope window or by using the mouse wheel. A double-click on the button resets the offset to the default of zero.

Figure 4.2 shows a zoomed-in version with a time offset of 180ms and a vertical zoom of +6dB applied to two signals to visually check them for identity. The test signal is routed to channels 1 and 2, the reference signal to channels 3 and 4 and the difference signal between them to channels 5 and 6. Hence the third lane shows the difference between the first two lanes.

Chapter 5

Regression Test

To test that the effects of JSFX-SoX really are bit-identical to SoX, a little test suite has been set up for checking DAW versus the command-line.

The suite assumes that command-line SoX is installed in the search path of your operating system.

If so, a simple batch script sets up raw audio test files and — externally via the command line — applies SoX effects to them producing audio result files. The parameters used are a bit exotic to ensure that algorithmic differences between SoX and JSFX-SoX will show up. The batch script can be found in the `test` subdirectory and is called `makeTestFiles.bat` (for Windows) or `makeTestFiles.sh` (for MacOS and Linux).

Since there are so many DAWs available, it is hard to provide a test project for each of those. The distribution just contains a Reaper project referencing those audio test files and result files in autonomous tracks (see figure 5.1). Adaption to other DAWs should be straightforward.

The JSFX-SoX effects are configured with the exactly the same parameters as given in the batch file and are correspondingly applied to the raw audio test files.

When subtracting the rendered audio in Reaper and the externally rendered audio from SoX, they (almost) cancel out. This can be checked by a spectrum analyser in the master channel, which is shown in figure 5.2. It shows a noise floor of typically less than -120dB.

Surprisingly the tracks do not completely cancel out, but this comes from rounding or precision errors — SoX often uses 32bit integer processing, while the JSFX-SoX always use double precision floating point processing — and also the 24 bit sample depth used in the FLAC files of the test suite; increasing that sample depth would even lead to less residual noise.



Figure 5.1: Regression Test Setup in Reaper




Figure 5.2: Example Noise Floor for Regression Test in Reaper

Bibliography

- [DBIQFILT] Wikipedia.
Digital Biquad Filter.
http://en.wikipedia.org/wiki/Digital_biquad_filter
- [FREEVERB] J. O. Smith.
Physical Audio Signal Processing.
W3K Publishing, 2010, ISBN 978-0-9745607-2-4.
<https://ccrma.stanford.edu/~jos/pasp/Freeverb.html>
- [JSFXSoX] Dr. Thomas Tensi.
SoX Plugins for Reaper DAW in JSFX Language.
<https://github.com/prof-spock/JSFX-SoX-Plugins>
- [REAPACK] Christian Fillion.
ReaPack: Package manager for REAPER.
<https://reapack.com>
- [REAPER] Cockos Incorporated.
Reaper Digital Audio Workstation.
<https://www.reaper.fm>
- [REAPLUGS] Cockos Incorporated.
ReaPlugs VST FX Suite.
<https://www.reaper.fm/reaplugs>
- [RBJFILT] R. Bristow-Johnson.
Cookbook formulae for audio EQ biquad filter coefficients.
<https://www.w3.org/2011/audio/audio-eq-cookbook.html>
- [SoXDOC] Chris Bagwell, Lance Norskog, Måns Rullgård et al.
SoX - SOund eXchange - Documentation.
<http://sox.sourceforge.net/Docs/Documentation>
- [SoXNG] Martin W. Guy.
SoX NG.
https://codeberg.org/sox_ng/sox_ng